

Jiawei ZHANG

jwz@uchicago.edu ◇ javyduck.github.io

(+1) 217-200-3511 ◇ Crerar 283, 5730 S Ellis Ave, Chicago, IL 60637

EDUCATION

University of Chicago

Ph.D. in Computer Science, 4.0/4.0

Advisor: *Prof. Bo Li*

Sept. 2024 – Aug. 2026

University of Illinois Urbana-Champaign

Ph.D. in Computer Science, 4.0/4.0

Advisor: *Prof. Bo Li*

Aug. 2023 – May. 2024

M.S. in Computer Science

May. 2023

Zhejiang University

Bachelor of Engineering

Hangzhou, China

Sept. 2017 – Jun. 2021

RESEARCH INTEREST:

I study **Safe AGI**, with a focus on **making LLMs and LLM agents reliable at scale**. My research runs as a closed loop: (i) **scalable red-teaming** to elicit realistic, long-horizon failures. I participated in internal red-teaming evaluations of **OpenAI o1**, **Google DeepMind**, and **ElevenLabs TTS** with **Virtue AI**. In parallel, I study **agentic reliability** for LLM agents, focusing on agent-specific vulnerabilities (e.g., prompt injection and malicious action execution) as well as **environment security construction** for agentic systems; (ii) **defense & alignment** that transforms safety requirements into actionable safeguards, spanning **certified robustness** for worst-case guarantees and **interpretability-guided defenses** that translate mechanistic insights into generalizable mitigations; and (iii) **better safety-aware reasoning for autonomous driving** by incorporating richer world knowledge via reinforcement learning for vision–language–action models.

Together, I aim to build **scalable red-teaming pipelines for LLM agents** and develop **practical patches** that improve safety awareness **while maintaining scalability**. Previously, I was an **Anthropic Fellow for AI Safety Research** from January to July 2026.

SELECTED PUBLICATION (* DENOTES CO-FIRST AUTHORSHIP)

- **Jiawei Zhang**, Andrew Estornell, David D. Baek, Bo Li, Xiaojun Xu. Any-Depth Alignment: Unlocking Innate Safety Alignment of LLMs to Any-Depth. *International Conference on Learning Representations (ICLR) 2026*. [arXiv]
- Zhaorun Chen*, Xun Liu*, Mintong Kang, **Jiawei Zhang**, Minzhou Pan, Shuang Yang, Bo Li. ARMs: Adaptive Red-Teaming Agent against Multimodal Models with Plug-and-Play Attacks. *International Conference on Learning Representations (ICLR) 2026*. [arXiv]
- Mintong Kang*, Zhaorun Chen*, Chejian Xu*, **Jiawei Zhang***, Chengquan Guo*, Minzhou Pan, Ivan Revilla, Yu Sun, Bo Li. GuardSet-X: Massive Multi-Domain Safety Policy-Grounded Guardrail Dataset. *NeurIPS 2025*. [arXiv]
- **Jiawei Zhang**, Shuang Yang, Bo Li. UDora: A Unified Red Teaming Framework Against LLM Agents by Dynamically Leveraging Their Own Reasoning. *International Conference on Machine Learning (ICML) 2025*. [arXiv]
- **Jiawei Zhang**, Xuan Yang, Taiqi Wang, Yu Yao, Aleksandr Petiushko, Bo Li. SafeAuto: Knowledge-Enhanced Safe Autonomous Driving with Multimodal Foundation Models. *International Conference on Machine Learning (ICML) 2025*. [arXiv]
- Chejian Xu*, **Jiawei Zhang***, Zhaorun Chen*, Chulin Xie*, Mintong Kang*, Zhuowen Yuan*, Chenhui Zhang, Lingzhi Yuan, Yi Zeng, Peiyang Xu, Chengquan Guo, Andy Zhou, ..., Zidi Xiong, Zinan Lin, Dan Hendrycks, Dawn Song, Bo Li. MMDT: Decoding the Trustworthiness and Safety of Multimodal Foundation Models. *International Conference on Learning Representations (ICLR) 2025*. [arXiv]
- **Jiawei Zhang**, Chejian Xu, Yu Gai, Freddy Lecue, Dawn Song, Bo Li. KnowHalu: Hallucination Detection via Multi-Form Knowledge Based Factual Checking. *ICLR 2025 Workshop on Foundation Models in the Wild*. [arXiv]
- Bowen Jin, Chulin Xie, **Jiawei Zhang**, Kashob Kumar Roy, Yu Zhang, Zheng Li, Ruirui Li, Xianfeng Tang, Suhang Wang, Yu Meng, Jiawei Han. Graph Chain-of-Thought: Augmenting Large Language Models by Reasoning on Graphs. *Findings of the Association for Computational Linguistics (ACL) 2024*. [arXiv]

- **Jiawei Zhang**, Chejian Xu, Bo Li. ChatScene: Knowledge-Enabled Safety-Critical Scenario Generation for Autonomous Vehicles. *Conference on Computer Vision and Pattern Recognition (CVPR) 2024*. [arXiv]
- **Jiawei Zhang**, Tianyu Pang, Chao Du, Yi Ren, Bo Li, Min Lin. MMCBench: Benchmarking Large Multimodal Models against Common Corruptions. [arXiv]
- **Jiawei Zhang**, Zhongzhu Chen, Huan Zhang, Chaowei Xiao, Bo Li. DiffSmooth: Certifiably Robust Learning via Diffusion Models and Local Smoothing. *32th USENIX Security Symposium 2023*. [arXiv]
- **Jiawei Zhang**, Linyi Li, Ce Zhang, Bo Li. CARE: Certifiably Robust Learning with Reasoning via Variational Inference. *IEEE Conference on Secure and Trustworthy Machine Learning (SatML) 2023*. [arXiv]
- Zhuolin Yang*, Zhikuan Zhao*, Boxin Wang, **Jiawei Zhang**, Linyi Li, Hengzhi Pei, Bojan Karlas, Ji Liu, Heng Guo, Ce Zhang, Bo Li. Improving Certified Robustness via Statistical Learning with Logical Reasoning. *Advances in Neural Information Processing Systems (NeurIPS) 2022*. [arXiv]
- Linyi Li, **Jiawei Zhang**, Tao Xie, Bo Li. Double Sampling Randomized Smoothing. *International Conference on Machine Learning (ICML) 2022*. [arXiv]
- **Jiawei Zhang***, Linyi Li*, Huichen Li, Xiaolu Zhang, Shuang Yang, Bo Li. Progressive-Scale Boundary Blackbox Attack via Projective Gradient Estimation. *International Conference on Machine Learning (ICML) 2021*. [arXiv]

INDUSTRY RESEARCH EXPERIENCE

OpenAI

Safety Research, Member of Technical Staff

Aug 2026 – Present

Anthropic

AI Safety Research Fellow

Jan 2026 – July 2026

- Automatic generation of environments and attacks for prompt injection evaluation at scale.

NVIDIA

ML Research Intern, Santa Clara

Oct 2025 – Jan 2026

Advised by Dr. Wenhao Ding, Dr. Yulong Cao, Dr. Boris Ivanovic, and Prof. Marco Pavone

- Explored improved reinforcement learning methods for reasoning to reduce minADE.
- Developed enhanced reasoning capabilities for autonomous driving by leveraging world models.

ByteDance Seed Research

Responsible AI Research Intern, San Jose

June 2025 – Sept 2025

Advised by Dr. Xiaojun Xu and Dr. Hang Li

- Developed a scalable, model-agnostic defense for LLMs based on safety-signature tokens, effective against deep prefill attacks with context lengths exceeding 3,000 tokens.
- Demonstrated robustness to in-the-wild jailbreaks (AutoDAN, PAIR) while maintaining 0% false positive rate on standard benchmarks (MMLU, MATH, BBH, AlphaEval).

Nuro AI

Machine Learning Research Intern (Pathfinder), Mountain View

May 2024 – Aug 2024

Advised by Dr. Aleksandr Petiushko

- Fine-tuned a multimodal LLM for video-conditioned autonomous driving to output high-level action plans and low-level control signals, unifying perception, reasoning, and control.
- Developed a multimodal retrieval-augmented training pipeline and a Markov Logic Network that encodes first-order traffic rules to verify and improve the safety of LLM-proposed actions.
- Proposed a position-dependent cross-entropy loss (PDCE) to stabilize and improve numeric control predictions when represented as text.

Sea AI Lab

Machine Learning Research Intern, Singapore

May 2023 – Aug 2023

Advised by Dr. Tianyu Pang & Dr. Chao Du

- Conducted evaluations on cross-modal models (e.g., Stable Diffusion, Whisper) to assess their consistency under a range of common data corruptions.
- Developed a rigorous benchmark for assessing the self-consistency of these models. The benchmark was designed to provide a comprehensive understanding of model behavior, incorporating a wide range of scenarios and inputs to measure their resilience and accuracy.

WORKSHOPS AND COMPETITIONS

- **CLAS 2024: The Competition for LLM and Agent Safety, NeurIPS 2024**
- **Secure and Safe Autonomous Driving (SSAD) Workshop and Challenge, CVPR 2023**

TEACHING

CS 307 - Modeling and Learning in Data Science (Spring 2022)

- Teaching Assistant with *Prof.* Bo Li and *Prof.* David Forsyth

SELECTED HONOR & AWARDS

Graduate Conference Travel Award	<i>March. 2023</i>
Accumulative Research Innovation & Entrepreneurship Index: 1st Department-wide	<i>June. 2021</i>
Meritorious Winner, MCM COMAP's Mathematical Contest in Modeling	<i>Apr. 2020</i>
First Prize, China Harbour Scholarship	<i>Jan. 2020</i>
First Prize, The Chinese Mathematics Competitions, Zhejiang Province	<i>Nov. 2018</i>

PROFESSIONAL SERVICE

Session Chair — ICLR 2025

Top Reviewer — NeurIPS 2025

Program Committee— ICML, NeurIPS, ICLR, CVPR, ECCV, AISTATS, AAAI, COLM, ACL, JIMR, etc.